

Decisions and randomness

Christos Dimitrakakis

March 27, 2025

Outline

Statistical Decision Theory

- Elementary Decision Theory

- Statistical Decision Theory

Gradient methods

- Gradients for optimisation

Statistical Decision Theory

Elementary Decision Theory

Statistical Decision Theory

Gradient methods

Gradients for optimisation

Preferences

Types of rewards

- ▶ For e.g. a student: Tickets to concerts.
- ▶ For e.g. an investor: A basket of stocks, bonds and currency.
- ▶ For everybody: Money.

Preferences among rewards

For any rewards $x, y \in R$, we either

- ▶ (a) Prefer x at least as much as y and write $x \succeq^* y$.
- ▶ (b) Prefer x not more than y and write $x \preceq^* y$.
- ▶ (c) Prefer x about the same as y and write $x \sim^* y$.
- ▶ (d) Similarly define \succ^* and \prec^*

Utility and Cost

Utility function

To make it easy, assign a utility $U(x)$ to every reward through a utility function $U : R \rightarrow \mathbb{R}$.

Utility-derived preferences

We prefer items with higher utility, i.e.

- ▶ (a) $U(x) \geq U(y) \Leftrightarrow x \succeq^* y$
- ▶ (b) $U(x) \leq U(y) \Leftrightarrow y \succeq^* x$

Cost

It is sometimes more convenient to define a cost function $C : R \rightarrow \mathbb{R}$ so that we prefer items with lower cost, i.e.

- ▶ $C(x) \geq C(y) \Leftrightarrow y \succeq^* x$

Random outcomes

Choosing among rewards

-[A] Bet 10 CHF on black -[B] Bet 10 CHF on 0 -[C] Bet nothing What is the reward here?

Choosing among trips

-[A] Taking the car to Zurich (50' without delays, 80' with delays) -[B] Taking the train to Zurich (60' without delays) What is the reward here?

Random rewards

- ▶ Each gamble gives us different rewards with different probabilities.
- ▶ These rewards are then **random**
- ▶ For simplicity, we assign a real-valued **utility** to outcomes. This is a **random variable**

Statistical Decision Theory

Elementary Decision Theory

Statistical Decision Theory

Gradient methods

Gradients for optimisation

Expected utility

Actions, outcomes and utility

In this setting, we obtain random outcomes that depend on our actions.

- ▶ Actions $a \in A$
- ▶ Outcomes $\omega \in \Omega$.
- ▶ Probability of outcomes $P(\omega \mid a)$
- ▶ Utility $U : \Omega \rightarrow \mathbb{R}$

Expected utility

The expected utility of an action is:

$$\mathbb{E}_P[U \mid a] = \sum_{\omega \in \Omega} U(\omega)P(\omega \mid a).$$

The expected utility hypothesis

We prefer a to a' if and only if

$$\mathbb{E}_P[U \mid a] \geq \mathbb{E}_P[U \mid a']$$

The St-Petersburg Paradox

The game

If you give me x CHF, then I promise to (a) Throw a fair coin until it comes heads. (b) If it does so after T throws, then I will give you 2^T CHF.



The question

- ▶ How much x are you willing to pay to play?
- ▶ Given that the expected amount of money is infinite, why are you only willing to pay a small x ?

Example: Betting

In this example, probabilities reflect actual randomness

Choice	Win Probability p	Payout w	Expected gain
Don't play	0	0	0
Black	18/37	2	
Red	18/37	2	
0	1/37	36	
1	1/37	36	



What are the expected gains for these bets?

Example: Route selection

- ▶ In this example, probabilities reflect subjective beliefs

Choice	Best time	Chance of delay	Delay amount	Expected time
Train	80	5%	5	
Car, route A	60	50%	30	
Car, route B	70	10%	10	

Example: Estimation

- In this example, probabilities are calculated starting from subjective beliefs

Mean-Square Estimation

If we want to guess $\hat{\theta}$, and we knew that $\theta \sim P$, then the guess

$$\hat{\theta} = \mathbb{E}_P(\theta) = \arg \min_{\hat{\theta}} \mathbb{E}_P[(\theta - \hat{\theta})^2]$$

minimises the squared error. This is because

$$\frac{d}{d\hat{\theta}} \mathbb{E}_P[(\theta - \hat{\theta})^2] = \frac{d}{d\hat{\theta}} \sum_{\omega} [\theta(\omega) - \hat{\theta}]^2 P(\omega) \quad (1)$$

$$= \sum_{\omega} \frac{d}{d\hat{\theta}} [\theta(\omega) - \hat{\theta}]^2 P(\omega) \quad (2)$$

$$= \sum_{\omega} 2[\theta(\omega) - \hat{\theta}](-1)P(\omega) = 2(\hat{\theta} - \mathbb{E}_P[\theta]). \quad (3)$$

Setting this to 0 gives $\hat{\theta} = \mathbb{E}_P[\theta]$

Example: Noisy optimisation

We wish to find the maximum of a function

$$f(x) \triangleq \mathbb{E}[g|x], \quad \mathbb{E}[g|x] = \int_{-\infty}^{\infty} g(\omega, x) p(\omega) d\omega \quad (4)$$

For this problem we need to use some more complex optimisation method, such as gradient methods

Statistical Decision Theory

Elementary Decision Theory

Statistical Decision Theory

Gradient methods

Gradients for optimisation

The gradient descent method: one dimension

- ▶ Function to minimise $f : \mathbb{R} \rightarrow \mathbb{R}$.
- ▶ Derivative $\frac{d}{d\theta} f(\beta)$

Gradient descent algorithm

- ▶ Input: initial value θ^0 , **learning rate** schedule α_t
- ▶ For $t = 1, \dots, T$
 - ▶ $\theta^{t+1} = \theta^t - \alpha_t \frac{d}{d\theta} f(\theta^t)$
- ▶ Return θ^T

Properties

- ▶ If $\sum_t \alpha_t = \infty$ and $\sum_t \alpha_t^2 < \infty$, it finds a local minimum θ^T , i.e. there is $\epsilon > 0$ so that

$$f(\theta^T) < f(\theta), \forall \theta : \|\theta^T - \theta\| < \epsilon.$$

Gradient methods for expected value

Estimate the expected value

$x_t \sim P$ with $\mathbb{E}_P[x_t] = \mu$.

Gradient methods for expected value

Estimate the expected value

$x_t \sim P$ with $\mathbb{E}_P[x_t] = \mu$.

Objective: mean squared error

Here $\ell(x, \theta) = (x - \theta)^2$.

$$\min_{\theta} \mathbb{E}_P[(x_t - \theta)^2].$$

Gradient methods for expected value

Estimate the expected value

$x_t \sim P$ with $\mathbb{E}_P[x_t] = \mu$.

Objective: mean squared error

Here $\ell(x, \theta) = (x - \theta)^2$.

$$\min_{\theta} \mathbb{E}_P[(x_t - \theta)^2].$$

Exact gradient update

If we know P , then we can calculate

$$\theta^{t+1} = \theta^t - \alpha_t \frac{d}{d\theta} \mathbb{E}_P[(x - \theta^t)^2] \quad (5)$$

$$\frac{d}{d\theta} \mathbb{E}_P[(x - \theta^t)^2] = 2 \mathbb{E}_P[x] - \theta^t \quad (6)$$

Gradient for mean estimation

- ▶ Let us show this in detail


$$\begin{aligned}\frac{d}{d\theta} \mathbb{E}_P[(x - \theta)^2] &= \int_{-\infty}^{\infty} dP(x) \frac{d}{d\theta} (x - \theta)^2 \\ &= \int_{-\infty}^{\infty} dP(x) 2(x - \theta) \\ &= 2 \mathbb{E}_P[x] - 2\theta.\end{aligned}$$

- ▶ If we set the derivative to zero, then we find the optimal solution:

$$\theta^* = \mathbb{E}_P[x]$$

- ▶ How can we do this if we only have data $x_t \sim P$?

Mean-squared error cost function



Here we see a plot of $\ell(\mu, \beta) = (\beta - \mu)^2$.

Stochastic gradient for mean estimation

Theorem (Sampling)

For any bounded random variable f ,

$$\mathbb{E}_P[f] = \int_{\mathcal{X}} dP(x) f(x) = \lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=1}^T f(x_t) = \mathbb{E}_P \left[\frac{1}{T} \sum_{t=1}^T f(x_t) \right], \quad x_t \sim P$$

Example (Sampling)

► If we sample x we approximate the gradient:

$$\frac{d}{d\theta} \mathbb{E}_P[(x - \theta)^2] = \int_{-\infty}^{\infty} dP(x) \frac{d}{d\theta} (x - \theta)^2 \approx \frac{1}{T} \sum_{t=1}^T \frac{d}{d\theta} (x_t - \theta)^2 = \frac{1}{T} \sum_{t=1}^T 2(x_t - \theta)$$

Stochastic gradient for mean estimation

Theorem (Sampling)

For any bounded random variable f ,

$$\mathbb{E}_P[f] = \int_{\mathcal{X}} dP(x) f(x) = \lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=1}^T f(x_t) = \mathbb{E}_P \left[\frac{1}{T} \sum_{t=1}^T f(x_t) \right], \quad x_t \sim P$$

Example (Sampling)

- If we sample x we approximate the gradient:

$$\frac{d}{d\theta} \mathbb{E}_P[(x - \theta)^2] = \int_{-\infty}^{\infty} dP(x) \frac{d}{d\theta} (x - \theta)^2 \approx \frac{1}{T} \sum_{t=1}^T \frac{d}{d\theta} (x_t - \theta)^2 = \frac{1}{T} \sum_{t=1}^T 2(x_t - \theta)$$

- If we update θ after each new sample x_t , we obtain:

$$\theta^{t+1} = \theta^t + 2\alpha_t(x_t - \theta^t)$$

The gradient method

- ▶ Function to minimise $f : \mathbb{R}^n \rightarrow \mathbb{R}$.
- ▶ Derivative $\nabla_{\theta} f(\theta) = \left(\frac{\partial f(\theta)}{\partial \theta_1}, \dots, \frac{\partial f(\theta)}{\partial \theta_n} \right)$, where $\frac{\partial f}{\partial \theta_n}$ denotes the **partial** derivative, i.e. varying one argument and keeping the others fixed.

Gradient descent algorithm

- ▶ Input: initial value θ^0 , learning rate schedule α_t
- ▶ For $t = 1, \dots, T$
 - ▶ $\theta^{t+1} = \theta^t - \alpha_t \nabla_{\theta} f(\theta^t)$
- ▶ Return θ^T

Properties

- ▶ If $\sum_t \alpha_t = \infty$ and $\sum_t \alpha_t^2 < \infty$, it finds a local minimum θ^T , i.e. there is $\epsilon > 0$ so that

$$f(\theta^T) < f(\theta), \forall \theta : \|\theta^T - \theta\| < \epsilon.$$

Stochastic gradient method

This is the same as the gradient method, but with added noise:

- ▶ $\theta^{t+1} = \theta^t - \alpha_t [\nabla_{\theta} f(\theta^t) + \omega_t]$
- ▶ $\mathbb{E}[\omega_t] = 0$ is sufficient for convergence.

Stochastic gradient method

This is the same as the gradient method, but with added noise:

- ▶ $\theta^{t+1} = \theta^t - \alpha_t [\nabla_{\theta} f(\theta^t) + \omega_t]$
- ▶ $\mathbb{E}[\omega_t] = 0$ is sufficient for convergence.

Example (When the cost is an expectation)

In machine learning, the cost is frequently an expectation of some function ℓ ,

$$f(\theta) = \int_{\mathcal{X}} dP(x) \ell(x, \theta)$$

This can be approximated with a sample

$$f(\theta) \approx \frac{1}{T} \sum_t \ell(x_t, \theta)$$

The same holds for the gradient:

$$\nabla_{\theta} f(\theta) = \int_{\mathcal{X}} dP(x) \nabla_{\theta} \ell(x, \theta) \approx \frac{1}{T} \sum_t \nabla_{\theta} \ell(x_t, \theta)$$