

Uninformed search

Christos Dimitrakakis

March 6, 2024

Outline

Shortest path algorithms

The shortest path problem

The shortest path problem

- ▶ Traversing arc $\langle x, y \rangle$ incurs **costs** $c(\langle x, y \rangle)$

The shortest path problem

- ▶ Traversing arc $\langle x, y \rangle$ incurs **costs** $c(\langle x, y \rangle)$
- ▶ Following a **path** h has a total cost $C(h) = \sum_{\langle x, y \rangle \in h} c(\langle x, y \rangle)$

The shortest path problem

- ▶ Traversing arc $\langle x, y \rangle$ incurs **costs** $c(\langle x, y \rangle)$
- ▶ Following a **path** h has a total cost $C(h) = \sum_{\langle x, y \rangle \in h} c(\langle x, y \rangle)$
- ▶ We can equivalently consider state-action **costs** $c(s, a)$.

The shortest path problem

- ▶ Traversing arc $\langle x, y \rangle$ incurs **costs** $c(\langle x, y \rangle)$
- ▶ Following a **path** h has a total cost $C(h) = \sum_{\langle x, y \rangle \in h} c(\langle x, y \rangle)$
- ▶ We can equivalently consider state-action **costs** $c(s, a)$.
- ▶ A policy π specifies a path x_1, \dots with $x_{k+1} = \tau(x_k, \pi(x_k))$

The shortest path problem

- ▶ Traversing arc $\langle x, y \rangle$ incurs **costs** $c(\langle x, y \rangle)$
- ▶ Following a **path** h has a total cost $C(h) = \sum_{\langle x, y \rangle \in h} c(\langle x, y \rangle)$
- ▶ We can equivalently consider state-action **costs** $c(s, a)$.
- ▶ A policy π specifies a path x_1, \dots with $x_{k+1} = \tau(x_k, \pi(x_k))$
- ▶ Following a **policy** π from state $x_1 = x$ has a total cost $C^\pi(x_1) = \sum_{k=1}^t c(x_k, \pi(x_k))$.

The shortest path problem

- ▶ Traversing arc $\langle x, y \rangle$ incurs **costs** $c(\langle x, y \rangle)$
- ▶ Following a **path** h has a total cost $C(h) = \sum_{\langle x, y \rangle \in h} c(\langle x, y \rangle)$
- ▶ We can equivalently consider state-action **costs** $c(s, a)$.
- ▶ A policy π specifies a path x_1, \dots with $x_{k+1} = \tau(x_k, \pi(x_k))$
- ▶ Following a **policy** π from state $x_1 = x$ has a total cost $C^\pi(x_1) = \sum_{k=1}^t c(x_k, \pi(x_k))$.

The shortest path problem

- ▶ Traversing arc $\langle x, y \rangle$ incurs **costs** $c(\langle x, y \rangle)$
- ▶ Following a **path** h has a total cost $C(h) = \sum_{\langle x, y \rangle \in h} c(\langle x, y \rangle)$
- ▶ We can equivalently consider state-action **costs** $c(s, a)$.
- ▶ A policy π specifies a path x_1, \dots with $x_{k+1} = \tau(x_k, \pi(x_k))$
- ▶ Following a **policy** π from state $x_1 = x$ has a total cost $C^\pi(x_1) = \sum_{k=1}^t c(x_k, \pi(x_k))$.

The shortest path problem

- ▶ Input: **start** nodes X and **goal** nodes Y and edge costs $c : A \rightarrow \mathbb{R}$.
- ▶ Output: Find a path h from X to Y so that $C(h) \leq C(h')$ for all h'

The shortest path problem

- ▶ Traversing arc $\langle x, y \rangle$ incurs **costs** $c(\langle x, y \rangle)$
- ▶ Following a **path** h has a total cost $C(h) = \sum_{\langle x, y \rangle \in h} c(\langle x, y \rangle)$
- ▶ We can equivalently consider state-action **costs** $c(s, a)$.
- ▶ A policy π specifies a path x_1, \dots with $x_{k+1} = \tau(x_k, \pi(x_k))$
- ▶ Following a **policy** π from state $x_1 = x$ has a total cost $C^\pi(x_1) = \sum_{k=1}^t c(x_k, \pi(x_k))$.

The shortest path problem

- ▶ Input: **start** nodes X and **goal** nodes Y and edge costs $c : A \rightarrow \mathbb{R}$.
- ▶ Output: Find a path h from X to Y so that $C(h) \leq C(h')$ for all h'

Notes

- ▶ If the path/policy does not reach a goal, the cost is infinite.
- ▶ We can maximise rewards instead of minimising costs.

Formalising the shortest path problem

The cost from state x of a policy that reaches a goal is

$$C^\pi(s) \triangleq \sum_{i=1}^{\infty} c[s_t, \pi(s_t)], \quad s_{t+1} = \tau[s_t, \pi(s_t)], \quad s_1 = s$$

where for every $s \in Y$, $c(s, a) = 0$ and $\tau(s, a) = s$ for all actions.

- We can calculate this recursively (from the goal state)

$$C^\pi(s) = \sum_{i=1}^{\infty} c[s_t, \pi(s_t)] \tag{1}$$

$$= c[s, \pi(s)] + \sum_{i=2}^{\infty} c[s_t, \pi(s_t)] \tag{2}$$

$$= c[s, \pi(s)] + C^\pi\{\tau[s, \pi(s)]\}. \tag{3}$$

- The same idea applies for the **shortest** path

$$C^*(s) \triangleq \min_{\pi} C^\pi(s) = \min_a \{c[s, a] + C^*[\tau(s, a)]\}. \tag{4}$$

The shortest path algorithm: backward search

Shortest path algorithm

Input: Goal states Y , starting state x .

Set $C(s) = 0$ for all states $s \in Y$, $F_0 = Y$.

for $t = 0, 1, \dots$ **do**

for $s' \in F_t$ **do**

$\pi(s) = \arg \min_a c(s, a) + C(\tau(s, a))$

$C(s) = \min_a c(s, a) + C(\tau(s, a))$

end for

$F_{t+1} = \text{parent}(F_t)$.

if $F_{t+1} = \emptyset$ or $x \in F_t$ **then**

return π, C

end if

end for

Algorithm idea

- ▶ Start from goal states
- ▶ Go back one step each time, adding the cost.
- ▶ Stop whenever there are no more states to go back to, or if we reach the start state.

Optimality proof

Theorem

$$C(s) = C^*(s)$$

Proof

- ▶ If $s \in Y$, then $C(s) = 0 = C^*(s)$.
- ▶ For any other s' , $s = \text{parent}(s')$: we will show that: if $C(s') \leq C^*(s')$ then $C(s) \leq C^*(s)$.

$$\begin{aligned} C(s) &= \min_a \{c(s, a) + C(\tau(s, a))\} && \text{(by definition)} \\ &\leq \min_a \{c(s, a) + C^*(\tau(s, a))\} && \text{(by induction)} \\ &\leq \min_a \left\{ c(s, a) + C^{\pi'}(\tau(s, a)) \right\}, \quad \forall \pi' && \text{(by optimality)} \\ &\leq C^\pi(s), \quad \forall \pi. \end{aligned}$$

For the optimal policy π^* , $C^{\pi^*}(s) = C^*(s)$, so $C(s) \leq C^*(s)$. Finally,

$$C^*(s) \leq C^\pi(s) = C(s) \geq C^*(s),$$

since $C^\pi(s) = C(s)$ for the policy returned by the algorithm.