

Informed search

Christos Dimitrakakis

March 8, 2024

Outline

The Shortest Path Problem

- The shortest path problem

- Heuristic Search

- Upper and lower bounds algorithms

General weight shortest path

- General weight shortest path

The shortest path problem

- ▶ Traversing arc $\langle x, y \rangle$ incurs **costs** $c(\langle x, y \rangle)$

The shortest path problem

- ▶ Traversing arc $\langle x, y \rangle$ incurs **costs** $c(\langle x, y \rangle)$
- ▶ Following a **path** h has a total cost $C(h) = \sum_{\langle x, y \rangle \in h} c(\langle x, y \rangle)$

The shortest path problem

- ▶ Traversing arc $\langle x, y \rangle$ incurs **costs** $c(\langle x, y \rangle)$
- ▶ Following a **path** h has a total cost $C(h) = \sum_{\langle x, y \rangle \in h} c(\langle x, y \rangle)$
- ▶ We can equivalently consider state-action **costs** $c(s, a)$.

The shortest path problem

- ▶ Traversing arc $\langle x, y \rangle$ incurs **costs** $c(\langle x, y \rangle)$
- ▶ Following a **path** h has a total cost $C(h) = \sum_{\langle x, y \rangle \in h} c(\langle x, y \rangle)$
- ▶ We can equivalently consider state-action **costs** $c(s, a)$.
- ▶ A policy π specifies a path x_1, \dots with $x_{k+1} = \tau(x_k, \pi(x_k))$

The shortest path problem

- ▶ Traversing arc $\langle x, y \rangle$ incurs **costs** $c(\langle x, y \rangle)$
- ▶ Following a **path** h has a total cost $C(h) = \sum_{\langle x, y \rangle \in h} c(\langle x, y \rangle)$
- ▶ We can equivalently consider state-action **costs** $c(s, a)$.
- ▶ A policy π specifies a path x_1, \dots with $x_{k+1} = \tau(x_k, \pi(x_k))$
- ▶ Following a **policy** π from state $x_1 = x$ has a total cost $C^\pi(x_1) = \sum_{k=1}^t c(x_k, \pi(x_k))$.

The shortest path problem

- ▶ Traversing arc $\langle x, y \rangle$ incurs **costs** $c(\langle x, y \rangle)$
- ▶ Following a **path** h has a total cost $C(h) = \sum_{\langle x, y \rangle \in h} c(\langle x, y \rangle)$
- ▶ We can equivalently consider state-action **costs** $c(s, a)$.
- ▶ A policy π specifies a path x_1, \dots with $x_{k+1} = \tau(x_k, \pi(x_k))$
- ▶ Following a **policy** π from state $x_1 = x$ has a total cost $C^\pi(x_1) = \sum_{k=1}^t c(x_k, \pi(x_k))$.

The shortest path problem

- ▶ Traversing arc $\langle x, y \rangle$ incurs **costs** $c(\langle x, y \rangle)$
- ▶ Following a **path** h has a total cost $C(h) = \sum_{\langle x, y \rangle \in h} c(\langle x, y \rangle)$
- ▶ We can equivalently consider state-action **costs** $c(s, a)$.
- ▶ A policy π specifies a path x_1, \dots with $x_{k+1} = \tau(x_k, \pi(x_k))$
- ▶ Following a **policy** π from state $x_1 = x$ has a total cost $C^\pi(x_1) = \sum_{k=1}^t c(x_k, \pi(x_k))$.

The shortest path problem

- ▶ Input: **start** nodes X and **goal** nodes Y and edge costs $c : A \rightarrow \mathbb{R}$.
- ▶ Output: Find a path h from X to Y so that $C(h) \leq C(h')$ for all h'

The shortest path problem

- ▶ Traversing arc $\langle x, y \rangle$ incurs **costs** $c(\langle x, y \rangle)$
- ▶ Following a **path** h has a total cost $C(h) = \sum_{\langle x, y \rangle \in h} c(\langle x, y \rangle)$
- ▶ We can equivalently consider state-action **costs** $c(s, a)$.
- ▶ A policy π specifies a path x_1, \dots with $x_{k+1} = \tau(x_k, \pi(x_k))$
- ▶ Following a **policy** π from state $x_1 = x$ has a total cost $C^\pi(x_1) = \sum_{k=1}^t c(x_k, \pi(x_k))$.

The shortest path problem

- ▶ Input: **start** nodes X and **goal** nodes Y and edge costs $c : A \rightarrow \mathbb{R}$.
- ▶ Output: Find a path h from X to Y so that $C(h) \leq C(h')$ for all h'

Notes

- ▶ If the path/policy does not reach a goal, the cost is infinite.
- ▶ We can maximise rewards instead of minimising costs.

Formalising the shortest path problem

The cost from state s of a policy that reaches a goal in Y is

$$C^\pi(s) \triangleq \sum_{i=1}^{\infty} c[s_t, \pi(s_t)], \quad s_{t+1} = \tau[s_t, \pi(s_t)], \quad s_1 = s$$

where for every $y \in Y$, $y(s, a) = 0$ and $\tau(y, a) = y$ for all actions.

- We can calculate this recursively (from the goal state)

$$C^\pi(s) = \sum_{i=1}^{\infty} c[s_t, \pi(s_t)] \tag{1}$$

$$= c[s, \pi(s)] + \sum_{i=2}^{\infty} c[s_t, \pi(s_t)] \tag{2}$$

$$= c[s, \pi(s)] + C^\pi\{\tau[s, \pi(s)]\}. \tag{3}$$

- The same idea applies for the **shortest** path

$$C^*(s) \triangleq \min_{\pi} C^\pi(s) = \min_a \{c[s, a] + C^*[\tau(s, a)]\}. \tag{4}$$

Dijkstra's shortest path algorithm: backward search

Shortest path algorithm

Input: Goal states Y , starting state x .

Set $C(s) = 0$ for all states $s \in Y$, $F_0 = Y$.

for $t = 0, 1, \dots$ **do**

for $s' \in F_t$ **do**

$\pi(s) = \arg \min_a c(s, a) + C(\tau(s, a))$

$C(s) = \min_a c(s, a) + C(\tau(s, a))$

end for

$F_{t+1} = \text{parent}(F_t)$.

if $F_{t+1} = \emptyset$ or $x \in F_t$ **then**

return π, C

end if

end for

Algorithm idea

- ▶ Start from goal states
- ▶ Go back one step each time, adding the cost.
- ▶ Stop whenever there are no more states to go back to, or if we reach the start state.

Optimality proof

Theorem

$$C(s) = C^*(s)$$

Proof

- ▶ If $s \in Y$, then $C(s) = 0 = C^*(s)$.
- ▶ For any other s' , $s = \text{parent}(s')$: we will show that: if $C(s') \leq C^*(s')$ then $C(s) \leq C^*(s)$.

$$\begin{aligned} C(s) &= \min_a \{c(s, a) + C(\tau(s, a))\} && \text{(by definition)} \\ &\leq \min_a \{c(s, a) + C^*(\tau(s, a))\} && \text{(by induction)} \\ &\leq \min_a \left\{ c(s, a) + C^{\pi'}(\tau(s, a)) \right\}, \quad \forall \pi' && \text{(by optimality)} \\ &\leq C^{\pi}(s), \quad \forall \pi. \end{aligned}$$

For the optimal policy π^* , $C^{\pi^*}(s) = C^*(s)$, so $C(s) \leq C^*(s)$. Finally,

$$C^*(s) \geq C^{\pi}(s) = C(s) \geq C^*(s),$$

since $C^{\pi}(s) = C(s)$ for the policy returned by the algorithm.

Partial graphs

- ▶ Why do we need search?
- ▶ We do not want to calculate on the whole graph
- ▶ We use **search** to find the shortest path more efficiently (perhaps).
- ▶ We denote the total cost of some path x_1, \dots, x_t as:

$$C(x_1, \dots, x_t)$$

- ▶ The remaining cost from x_t to the goal using some policy π as

$$C^\pi(x_t)$$

Generic search

We define heuristic search in the context of shortest-path problems.

We now consider a general method for searching a node in the frontier.

input $G = \langle N, E \rangle$: Graph.

input $f : N \rightarrow \mathbb{R}$: evaluation function.

input x : Start node

function Heuristic Search(G, x, h)

$S' = \emptyset$: Nodes searched.

$F = \{x\}$. Initialise the frontier

$c_x = 0$. Initialise the cost of node x

while $F \neq \emptyset$ **do**

$n = \arg \min_{i \in F} f(i)$. Select "best" node.

$F = F \setminus \{n\}$. Remove n from the frontier.

if $n \notin S'$ **then**

$B = \text{child}(n) \setminus S'$. Get the set of unsearched children of n .

$\forall b \in B, b_i = c_n + c(n, b)$. Calculate the total cost to each child b .

$S' = S' \cup \{n\}$. Add n to the list of searched nodes.

$F = F \cup B$. Add n 's children to the frontier.

end if

end while

A* search

We now consider a general method for searching a node in the frontier.

input $G = \langle N, E \rangle$: Graph.

input $h : N \rightarrow \mathbb{R}$: heuristic function.

input x : Start node

function A-Star(G, x, h)

$S' = \emptyset$: Nodes searched.

$F = \{x\}$. Initialise the frontier

$c_x = 0$. Initialise the cost of node x

while $F \neq \emptyset$ **do**

$n = \arg \min_{i \in F} c_i + h(i)$. Select minimum cost + heuristic node.

$F = F \setminus \{n\}$. Remove n from the frontier.

if $n \notin S'$ **then**

$B = \text{child}(n) \setminus S'$. Get the set of unsearched children of n .

$\forall b \in B, b_i = c_n + c(n, b)$. Calculate the total cost to each child b .

$S' = S' \cup \{n\}$. Add n to the list of searched nodes.

$F = F \cup B$. Add n 's children to the frontier.

end if

end while

► You can see that $h = 0$ corresponds to minimum-cost search.

Admissible heuristics

- ▶ If h is arbitrary, then the search can fail.
- ▶ We need h to be admissible. In particular,

$$C^*(n) \geq h(n).$$

Admissibility of A^*

Theorem

A^* returns an optimal solution if

- ▶ The graph has a bounded branching factor.
- ▶ All costs are greater than $\epsilon > 0$
- ▶ The heuristic is admissible, i.e. $0 \leq h(n) \leq C^*(n)$ for all $n \in N$.

Proof

- ▶ **Existence.** There is a finite number of paths that will be explored, as the longest possible path to a goal is $C^*(0)/\epsilon$ and the branching factor is bounded.
- ▶ **Optimality.** The proof is by contradiction. Let us assume that A^* finds some $\pi \neq \pi^*$ so that $C(\pi) > C(\pi^*)$. That means that at some node n on the path there is an action a^* on the optimal policy, but we keep expanding the path x_1, x_2, \dots of π . However, since $C(\pi) > C(\pi^*)$ there must be some t such that $C(n, x_1, \dots, x_t) > C^{\pi^*}(n)$. But then, to expand π requires that $C(n, x_1, \dots, x_t) + h(x') < h(x) \leq C^{\pi^*}(n)$.

Calculating Upper and Lower Bounds

Starting from a set of leaf nodes S_0

Upper bound $U(s) \geq C^*(s)$ for $s \in S_0$

Setting $U(0) \geq C^*(0)$ and recursing:

$$U(s) = \min_{a \in A_s} c(s, a) + U[\tau(s, a)]$$

By induction, we can prove that this is an upper bound on C^* :

$$U(s) = \min_{a \in A_s} c(s, a) + U[\tau(s, a)] \geq \min_{a \in A_s} c(s, a) + C^*[\tau(s, a)] = C^*(s).$$

Lower bound $L(s) \leq C^*(s)$ for $s \in S_0$

$$L(s) = \min_{a \in A_s} c(s, a) + L[\tau(s, a)]$$

Similarly, we can prove that it is a lower bound:

$$L(s) = \min_{a \in A_s} c(s, a) + L[\tau(s, a)] \leq \min_{a \in A_s} c(s, a) + C^*[\tau(s, a)] = C^*(s)$$

Branch and bound

The algorithm is rather simple to describe in words.

- ▶ [1] Set $s = 0$.
- ▶ [1.1] Select action a^* minimising $c(s, a) + L(\tau(s, a))$.
- ▶ [1.2] Discard subtrees (s, a) for which $c(s, a) + L(\tau(s, a)) \geq c(s, a^*) + L(\tau(s, a^*))$.
- ▶ [1.3] Proceed to $s = \tau(s, a)$ and go to 1.1. unless we are at a leaf.
- ▶ [2] Expand the leaf node, and generate new leaf nodes with corresponding upper and lower bounds.
- ▶ [3] Calculate L, S for the corresponding subtree.
- ▶ [4] Go to 1.

General weight shortest path

- ▶ In this problem, actions can have positive or negative costs.
- ▶ Negative edges generate problems if we have cycles
- ▶ However, the basic algorithmic idea is again Dynamic Programming

Bellman-Ford Algorithm

In state-action notation, the algorithm is simply

- ▶ $C_0(0) = 0$, $C_i(0) = \infty$ for all $i \neq 0$.
- ▶ For $k \in 1, \dots, |S|$:

$$C_k(s) = \min_a c(s, a) + C_{k-1}(\tau(s, a))$$

Bellman-Ford Algorithm

```
 $C(0) = 0. C(i) = \infty, \text{ for } i \neq 0.$   
for  $i \in 1, \dots, |N| - 1$  do  
  for all edges  $(i, j)$  do  
    if  $C(i) + c(i, j) < C(j)$  then  
       $c(j) = C(i) + c(i, j)$   
    end if  
  end for  
end for  
for all edges  $(i, j)$  do  
  if  $C(i) + c(i, j) < C(j)$  then  
    error "Negative cycle"  
  end if  
end for
```

- ▶ Succinctly, the algorithm is just like Dijkstra, but it ensures it goes at most $|N| - 1$ times through all vertices, and has a sanity check as no more updates should be possible at the end.
- ▶ Instead of keeping a track of explored nodes, it uses the fact that C is initialised to infinity.